

# Optimizing Multi-Intersection Traffic Signal Control Using Deep Reinforcement Learning

Abdul Basit & M. Abdullah

LUMS

December 18, 2025

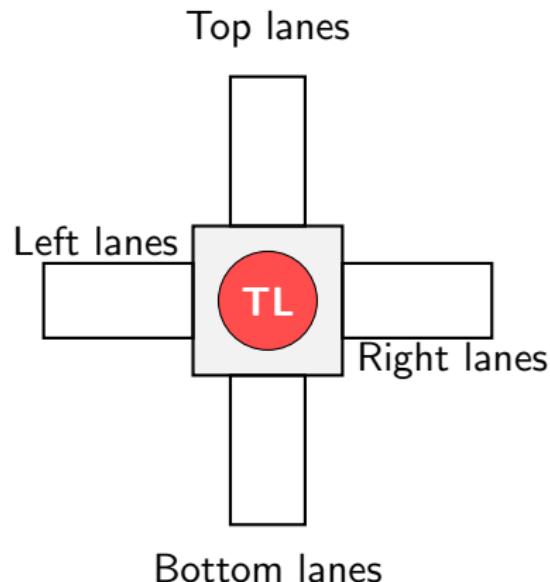
# Traffic Signal Control as an MDP Problem

- **Problem Statement:** Optimize traffic signal timing to minimize congestion and maximize throughput
- **MDP Components:**
  - States (S): Traffic conditions at intersections
  - Actions (A): Signal phase selections
  - Transitions (P): Traffic flow dynamics (SUMO simulator)
  - Rewards (R): Negative congestion + positive throughput
- **Objective:** Find policy  $\pi^*$  that maximizes expected cumulative reward

$$\pi^* = \arg \max_{\pi} \mathbb{E} \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) \mid \pi \right]$$

# Single Intersection MDP Setup

- **Environment:** One traffic light controlling 4 approaches (N, E, S, W)
- **Simulation:** SUMO microscopic traffic simulator with TraCI interface
- **Decision Frequency:** Every 5 seconds (action execution time)
- **Episode Length:** 3600 seconds (1 hour simulation)
- **Traffic Light Program:** 8 phases (4 green + transitions)
- **Agent Control:** Selects from green phases (0, 2, 4, 6)
- **Safe Transitions:** Environment handles yellow → red → green automatically



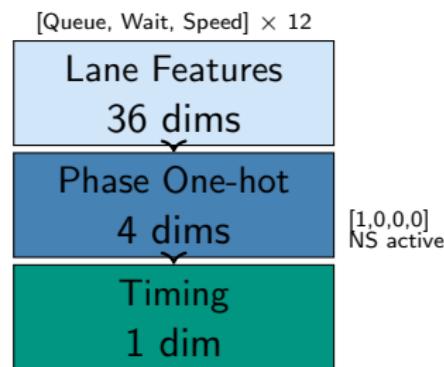
SUMO Phase Cycle: 0 → 1 → 2  
→ 3 → 4 → 5 → 6 → 7 → 0

# Single Intersection State Space (41 Dimensions)

- **State Vector Structure:** [lane\_features, phase\_encoding, timing]

- **Lane Features (36 dimensions):**

- 12 incoming lanes  $\times$  3 metrics each
- **Queue Length:** Number of halted vehicles per lane
- **Waiting Time:** Sum of waiting times for vehicles in lane
- **Speed:** Average normalized speed (0-1) / lane



- **Phase Encoding (4 dimensions):**

- One-hot vector for current green phase
- [1,0,0,0] = Phase 0 (NS through/right)
- [0,1,0,0] = Phase 2 (NS left)
- [0,0,1,0] = Phase 4 (EW through/right)
- [0,0,0,1] = Phase 6 (EW left)

- **Timing Feature (1 dimension):**

- Time since last phase change (normalized 0-1)

Total:  $36 + 4 + 1 = 41$  dimensions

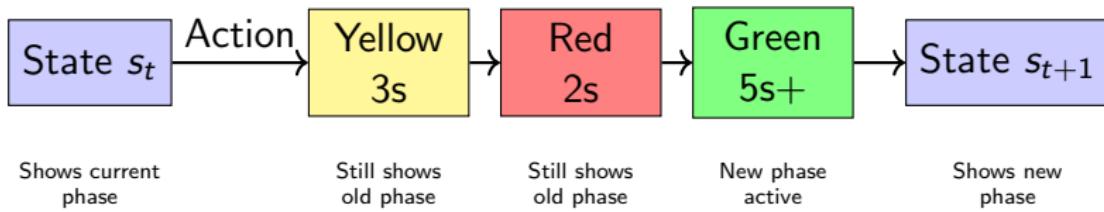
# Single Intersection State Space - Properties

## ■ State Update Timing:

- State reflects conditions **before** action execution
- During transitions: shows **previous** green phase
- Updates only after reaching new green phase

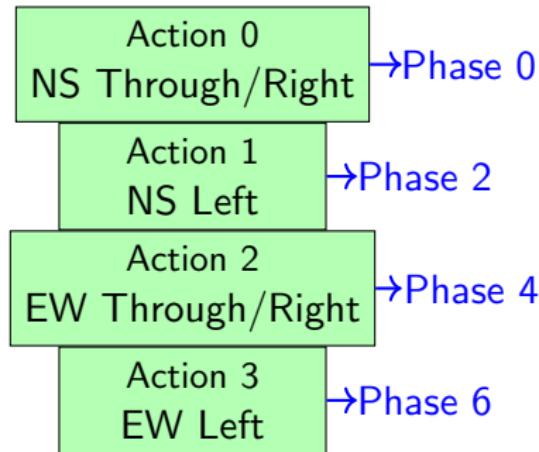
## ■ Why This State Representation?

- Captures local traffic congestion patterns
- Phase information enables temporal learning
- Normalized features handle varying traffic volumes
- Queue lengths indicate immediate congestion
- Waiting times measure cumulative delay
- Speed indicates flow quality



# Single Intersection Action Space (4 Actions)

- **Action Definition:** Direct selection of green traffic phases. Each state corresponds to SUMO phase with 12-character string
- **Action 0:** Phase 0 - NS through/right green (SUMO phase 0)
- **Action 1:** Phase 2 - NS left green (SUMO phase 2)
- **Action 2:** Phase 4 - EW through/right green (SUMO phase 4)
- **Action 3:** Phase 6 - EW left green (SUMO phase 6)



# Single Intersection Action Execution

## ■ Action Execution Process:

- ① Agent selects action (0, 1, 2, or 3)
- ② Environment checks:  $\text{action} \neq \text{current\_phase}$ ?
- ③ If different phase:
  - Set yellow phase for 3 seconds
  - Set red clearance phase for 2 seconds
  - Set new green phase
  - Update current\_phase and reset timer
- ④ If same phase: Continue current green phase
- ⑤ Execute simulation for 5 seconds with active phase
- ⑥ Compute reward and update state

## ■ Safe Transition Details:

- Yellow phase =  $\text{current\_phase} + 1$  (always)
- Red clearance uses all-red pattern phase
- New green =  $\text{green\_phases}[\text{action}]$
- All transitions use SUMO's built-in phase program

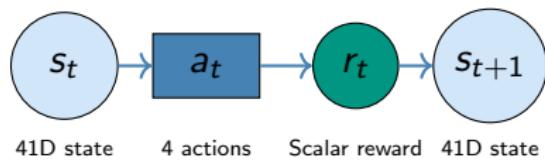
# Single Intersection MDP Components

- **States (S):** 41-dimensional continuous vectors
  - Lane features + phase encoding + timing
- **Actions (A):** 4 discrete phase selections  $\{0,1,2,3\}$
- **Transitions (P):** Deterministic given SUMO physics
  - $s_{t+1} = \text{SUMO\_dynamics}(s_t, a_t, \text{traffic\_arrivals})$
- **Rewards (R):** Traffic efficiency metric
  - Penalizes congestion, rewards throughput
- **Discount Factor ( $\gamma$ ):** 0.95
- **Horizon:** Finite episodes (3600s)

## Reward Function

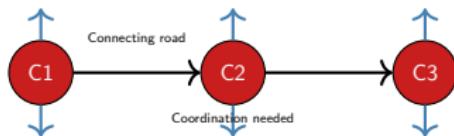
$$R = -0.5 \cdot W_t - Q_t + 500 \cdot T_t - P_t$$

- $W_t$  = Total waiting time
- $Q_t$  = Total queue length
- $T_t$  = Throughput
- $P_t$  = Congestion penalty



# Multi-Intersection MDP Setup (3 Intersections)

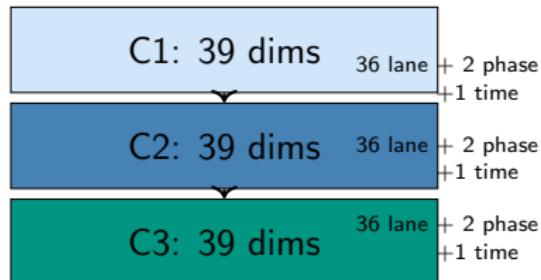
- **Environment:** 3 connected intersections (C1, C2, C3) in network
- **Coordination Challenge:** Joint control of correlated traffic flows
- **Network Topology:** Linear arrangement with connecting roads
- **Simulation:** Same SUMO setup with multi-TraCI connections
- **Decision Frequency:** Every 5 seconds (synchronized across intersections)
- **Episode Length:** 3600 seconds
- **Traffic Lights:** Each with 4-phase program (2 controllable greens)
- **Single Agent:** Controls entire network for coordinated optimization



Each intersection:  
12 lanes, 4 phases,  
2 controllable greens

# Multi-Intersection State Space (117 Dimensions)

- **State Structure:** Concatenated per-intersection features
- **Formula:** State = [C1\_features, C2\_features, C3\_features]
- **Per-Intersection Features (39 dimensions each):**
  - **Lane Features (36 dims):** 12 lanes × 3 metrics
    - Queue length, waiting time, normalized speed per lane
  - **Phase Encoding (2 dims):** One-hot for current green phase
    - $[1,0] = \text{NS green (Phase 0)}$ ,  $[0,1] = \text{EW green (Phase 2)}$
  - **Timing (1 dim):** Time since last phase change

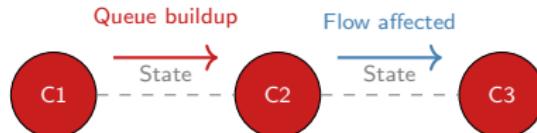


Total: 117 dimensions

- **Total Dimensions:** 3 intersections × 39 features = 117

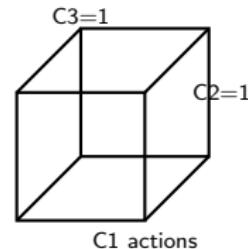
# Multi-Intersection State Space - Properties

- **Interdependence:** States capture cross-intersection traffic patterns
  - Queue buildup at C1 affects C2's incoming traffic
  - Phase coordination enables green wave effects
- **Phase Coordination:** Agent sees all intersection phases simultaneously
  - Can learn to synchronize compatible phases
  - Prevents conflicting traffic flows between intersections
- **Queue Propagation:** Downstream queues affect upstream decisions
  - Agent learns network-wide congestion dynamics
  - Enables proactive traffic management
- **Update Timing:** All intersections update states simultaneously
  - Synchronized across the entire network
  - Consistent state representation for coordinated actions



# Multi-Intersection Action Space (8 Joint Actions)

- **Individual Actions:** Each intersection has 2 actions
  - Action 0: NS green, Action 1: EW green
- **Joint Actions:** Cartesian product of individual actions
- **Action Space Size:**  $2^3 = 8$  joint actions
- **Encoding Scheme:**
  - Joint action =  $C_1\_action + 2 \times C_2\_action + 4 \times C_3\_action$
  - Example:  $[1,0,1] \rightarrow 1 + 0 + 4 =$  Joint Action 5
- **Decoding:** Extract per-intersection actions from joint index



$$2 \times 2 \times 2 = 8 \text{ joint actions}$$

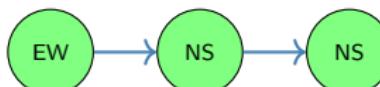
# Multi-Intersection Action Space - Interpretations

## ■ Joint Action Meanings:

- ① **Action 0:** [NS, NS, NS] - All intersections in north-south green
- ② **Action 1:** [EW, NS, NS] - C1 east-west, others north-south
- ③ **Action 2:** [NS, EW, NS] - C2 east-west, others north-south
- ④ **Action 3:** [EW, EW, NS] - C1 & C2 east-west, C3 north-south
- ⑤ **Action 4:** [NS, NS, EW] - C3 east-west, others north-south
- ⑥ **Action 5:** [EW, NS, EW] - C1 & C3 east-west, C2 north-south
- ⑦ **Action 6:** [NS, EW, EW] - C2 & C3 east-west, C1 north-south
- ⑧ **Action 7:** [EW, EW, EW] - All intersections in east-west green

## ■ Execution Details:

- All intersections transition simultaneously
- Each follows: green  $\rightarrow$  yellow(3s)  $\rightarrow$  red(2s)  $\rightarrow$  new\_green
- Independent transitions but coordinated timing



Joint Action 1: [EW, NS, NS]

C1: EW green    C2,C3: NS green

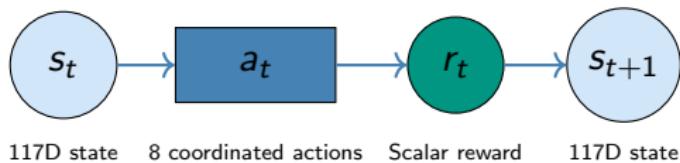
# Multi-Intersection Network MDP

- **States (S):** 117-dimensional continuous vectors
- **Actions (A):** 8 discrete joint phase selections  
 $\{0,1,2,3,4,5,6,7\}$
- **Transitions (P):** Network traffic dynamics
  - $s_{t+1} = \text{SUMO\_network\_dynamics}(s_t, a_t, \text{traffic\_patterns})$
- **Rewards (R):** Network-wide efficiency
  - Penalizes congestion, rewards throughput
- **Discount Factor ( $\gamma$ ):** 0.95
- **Horizon:** 3600s episodes

## Network Reward Function

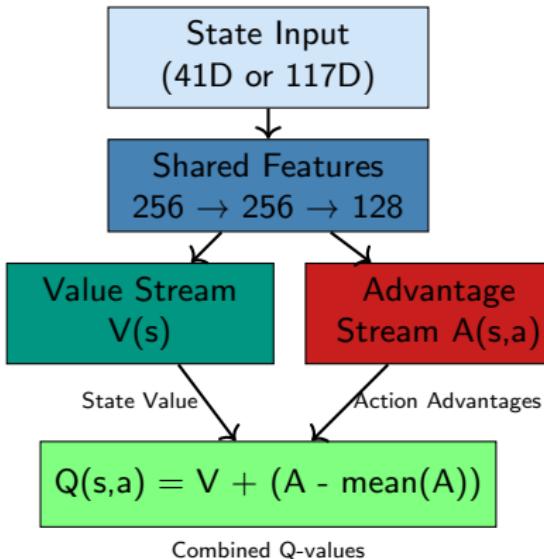
$$R = -0.5 \cdot W_t - Q_t + 100 \cdot T_t - P_t$$

- $W_t$  = total waiting time in the network
- $Q_t$  = total queue length
- $T_t$  = network throughput



# Dueling Double Deep Q-Network Architecture

- **Core Algorithm:** Dueling Double DQN combines two key innovations
- **Dueling Architecture:** Separates value and advantage estimation
- **Double DQN:** Reduces overestimation bias in Q-value estimation
- **Why for Traffic Control?**
  - High-dimensional state spaces (41D/117D)
  - Complex action coordination (4/8 actions)
  - Need for sample-efficient learning
  - Stability requirements for real-time control
- **Implementation:** Both enabled by default in my codebase



- **Traditional DQN:** Directly predicts  $Q(s, a)$ 
  - No separation of state value vs. action advantage
  - Inefficient learning of state values
- **Dueling DQN:** Separates value and advantage
  - **Value Stream:**  $V(s)$  – value of being in state  $s$
  - **Advantage Stream:**  $A(s, a)$  – advantage of action  $a$  in state  $s$
  - **Combination:**  $Q(s, a) = V(s) + (A(s, a) - \bar{A}(s))$

## Mathematical Formulation

$$Q(s, a) = V(s) + [A(s, a) - \bar{A}(s)]$$

$$\text{where: } \bar{A}(s) = \frac{1}{|A|} \sum_{a'} A(s, a')$$

## ■ Why subtract the mean?

- Ensures advantage has zero mean (identifiability)
- Makes  $V(s)$  the true state value
- Prevents gradient vanishing

# Advantages of Dueling DQN Architecture

- **Sample Efficiency:** Value stream learns state values; advantage stream learns action preferences. Fewer samples needed.
- **Better Generalization:** State values learned independently of actions. Can handle unseen action combinations, useful for traffic control.
- **Stable Learning:** Separates state and action learning, reducing interference and improving gradient flow.
- **Interpretability:** Analyze learned state values and action advantages for understanding policies.

## Why Perfect for Traffic Control?

State values matter regardless of actions; action advantages depend on traffic patterns. Efficient learning and coordination across intersections.

# Understanding Double DQN Algorithm

## ■ Problem: Overestimation Bias in DQN

- Same network selects and evaluates actions
- Inflated Q-values → suboptimal policies

## ■ Double DQN Solution

- **Policy Network:** selects best action
- **Target Network:** evaluates that action
- Decouples selection and evaluation

## Comparison

Standard DQN:

$$Q_{\theta^-}(s', \arg \max_{a'} Q_{\theta^-}(s', a'))$$

Double DQN:

$$Q_{\theta^-}(s', \arg \max_{a'} Q_{\theta}(s', a'))$$

## ■ Implementation

- Online network:  $a' = \arg \max_{a'} Q_{\theta}(s', a')$
- Target network:  $Q_{\theta^-}(s', a')$

# Advantages of Double DQN Algorithm

- **Reduces Overestimation Bias:**
  - More accurate Q-values → better policies
  - Important in sparse reward environments
- **More Stable Training:**
  - Less local optima, better convergence
  - Reduces noise impact
- **Better Exploration:**
  - Accurate Q-values → smarter exploration
  - $\epsilon$ -greedy more effective
- **Improved Performance:**
  - Outperforms standard DQN
  - Crucial for complex tasks like traffic control

## Why Critical for Traffic Control?

- Accurate Q-values for safety-critical decisions
- Overestimation could lead to dangerous traffic patterns
- Ensures stability for real-time control

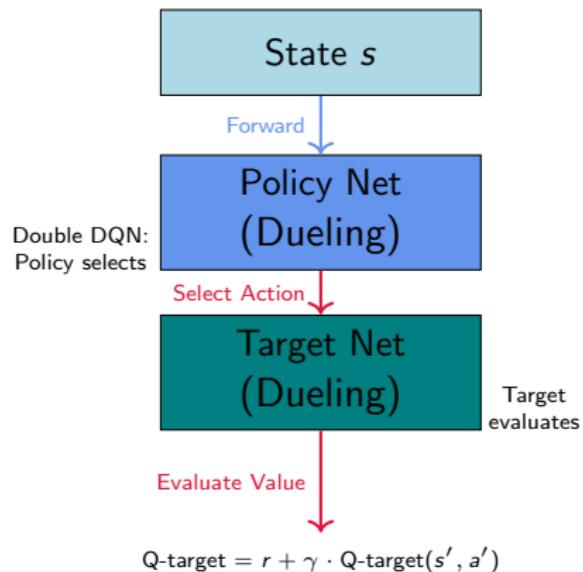
# Dueling Double DQN: Best of Both Worlds

## ■ Combined Architecture:

- Dueling for efficient value learning
- Double for accurate value estimation
- Target networks for stability
- Experience replay for sample efficiency

## ■ Training Process:

- ① Sample batch from experience replay
- ② Use policy network to select actions (Double DQN)
- ③ Use target network to evaluate values (Double DQN)
- ④ Update dueling network parameters
- ⑤ Periodically update target network



# Implementation Specifications

## Network Topology (Dueling)

- **Input Layer:**

- 41D (Single) or 117D (Multi)

- **Shared Feature Extractor:**

- FC Layers:  
[256 → 256 → 128]
  - Activation: ReLU

- **Dueling Streams (Split):**

- **Value  $V(s)$ :**  
128 → 128 → 1
  - **Advantage  $A(s, a)$ :**  
128 → 128 →  $|A|$

- **Output:**  $|A| = 4$  (Single) or 8 (Multi)

## Training Configuration

- **Optimization:**

- Optimizer: Adam
  - Learning Rate:  $1 \times 10^{-3}$
  - Gradient Clip: Max norm 10.0

- **Loss Function:**

- Huber Loss (Smooth L1)
  - Robust to outliers/noise

- **Stability Mechanisms:**

- **Double DQN:** Decoupled selection/evaluation
  - **Target Update:** Hard copy every 1000 steps

**Aggregation Formula:**  $Q(s, a) = V(s) + \left( A(s, a) - \frac{1}{|A|} \sum_{a'} A(s, a') \right)$

# Hyperparameter Configuration

## Optimization Dynamics

- **Learning Rate ( $\alpha$ ):**  $1 \times 10^{-3}$ 
  - Optimizer: Adam (Adaptive moments)
  - Rationale: Balance speed vs. stability
- **Discount Factor ( $\gamma$ ):** 0.99
  - Prioritizes long-term flow over immediate rewards
- **Experience Replay:**
  - **Batch Size:** 32 (Parallel processing)
  - **Buffer Size:** 10,000 transitions

## Stability & Exploration

- **$\epsilon$ -Greedy Schedule:**
  - Range: 1.0 → 0.01
  - Decay: 0.995 per episode
  - Rationale: High initial exploration
- **Network Stability:**
  - **Target Update:** Every 1000 steps
  - **Gradient Clip:** Max norm 10.0
- **Hidden Layers:**
  - Specs: [256, 256, 128]
  - Rationale: Capacity for state complexity without overfitting

# Why Dueling Double DQN for Traffic Signal Control?

## ■ High-Dimensional State Spaces:

- Single intersection: 41 dimensions
- Multi-intersection: 117 dimensions
- Dueling efficiently learns state values
- Double DQN provides accurate estimates

## ■ Complex Action Spaces:

- Single: 4 discrete actions
- Multi: 8 joint actions ( $2^3$  combinations)
- Dueling separates value vs advantage learning
- Better generalization across action combinations

## ■ Safety-Critical Application:

- Traffic control affects public safety
- Double DQN reduces overestimation risks
- Stable learning prevents dangerous policies
- Reliable performance in real-time settings

## ■ Sample Efficiency Requirements:

- SUMO simulations are computationally expensive
- Dueling architecture learns more efficiently
- Experience replay maximizes sample utilization

# SUMO Simulation Environment

## Simulator Platform

**SUMO** (Simulation of Urban MObility)

- Open-source microscopic traffic simulator
- Developed by DLR (German Aerospace Center)
- Industry-standard for traffic research

## TraCI Interface

- Real-time control via Python API
- Query traffic state every simulation step
- Set traffic light phases dynamically
- Retrieve vehicle positions, speeds, queues

# Network Topology & Infrastructure

## Road Network Design

- **Multi-Intersection Corridor:**
  - 3 signalized intersections (C1, C2, C3)
  - Total corridor length: 1000 meters
- **Each Intersection:**
  - 4 approaches (N, S, E, W)
  - 3 lanes per approach (12 lanes total)
- **Road Specifications:**
  - Speed limit: 50 km/h (13.89 m/s)
  - N/S exit distances: 200 meters

## Traffic Light Configuration

- **Phase Structure:**
  - Phase 0: North-South green
  - Phase 1: North-South yellow(3s)
  - Phase 2: East-West green
  - Phase 3: East-West yellow(3s)
- **Timing Constraints:**
  - Yellow duration: 3s (fixed)
  - Red clearance: 2s
  - Minimum green: 5s
  - Maximum green: 60s
- **Agent Control:**
  - Selects green phases only
  - Decision frequency: Every 5 seconds

# Traffic Generation & Demand Patterns

## Vehicle Characteristics

### ■ Vehicle Types:

- Standard cars: 90% of traffic
- Buses: 10% (longer, slower)

### ■ Car Parameters:

- Length: 5.0 meters
- Max speed: 60 km/h (16.67 m/s)
- Acceleration: 2.6 m/s<sup>2</sup>
- Deceleration: 4.5 m/s<sup>2</sup>

### ■ Bus Parameters:

- Length: 12.0 meters
- Max speed: 50 km/h (13.89 m/s)
- Acceleration: 1.2 m/s<sup>2</sup>
- Deceleration: 4.0 m/s<sup>2</sup>

## Demand Profiles

### ■ Time-Varying Traffic:

- Period 1 (0-900s): Low traffic
- Period 2 (900-1800s): Medium traffic
- Period 3 (1800-2700s): Peak traffic
- Period 4 (2700-3600s): Declining

### ■ Flow Probabilities:

- Low: 0.02 vehicles/second
- Medium: 0.03 vehicles/second
- Peak: 0.04 vehicles/second

# Simulation & Episode Configuration

## Episode Structure

- **Duration:** 3600 sec (1 h)
  - Captures full demand cycle
  - Includes congestion buildup/dissipation
- **Decision Frequency:** Every 5 sec
  - 720 decision steps per episode
  - Balance responsiveness vs stability
  - Matches practical deployment constraints
- **Training Episodes:**
  - Range: 200-500 episodes
  - Varies by hyperparameter configuration

## Metrics & Evaluation

- **Primary Metrics:**
  - **Waiting time:** Total time vehicles spend halted (seconds)
  - **Queue length:** Number of stopped vehicles
  - **Throughput:** Vehicles completing journey

# Systematic Hyperparameter Ablation Study

## Fixed Parameters

- Architecture: [256, 256, 128]
- Target network update: every 1000 steps
- Optimizer: Adam + gradient clipping (10.0)
- Loss: Smooth L1 (Huber)
- Episode length: 3600s (720 decisions)
- $\epsilon$ -greedy: 1.0 → 0.01

## Parameters Varied

- **Learning rate ( $\alpha$ ):** 0.0001 – 0.002 (stability vs convergence speed)
- **Discount factor ( $\gamma$ ):** 0.95 – 0.995 (short-term vs long-term planning)
- **Epsilon decay:** 0.98 – 0.999 (exploration-exploitation balance)
- **Batch size:** 64, 96, 128 (gradient estimate stability)
- **Replay buffer size:** 50k, 75k, 100k (experience diversity vs staleness)

## 12 Experimental Configurations

| Run | LR     | $\gamma$     | Batch     | $\epsilon$ Decay | Buf.        | Eps |
|-----|--------|--------------|-----------|------------------|-------------|-----|
| 01  | 0.0005 | 0.99         | 64        | 0.995            | 50k         | 500 |
| 02  | 0.0001 | 0.99         | 64        | 0.995            | 50k         | 500 |
| 03  | 0.001  | 0.99         | 128       | 0.995            | 50k         | 500 |
| 04  | 0.002  | 0.99         | 128       | 0.995            | 50k         | 500 |
| 05  | 0.0005 | 0.99         | 64        | <b>0.999</b>     | 50k         | 500 |
| 06  | 0.0005 | 0.99         | 64        | <b>0.99</b>      | 50k         | 500 |
| 07  | 0.0005 | 0.99         | 64        | <b>0.98</b>      | 50k         | 200 |
| 08  | 0.0005 | <b>0.95</b>  | 64        | 0.995            | 50k         | 200 |
| 09  | 0.0005 | <b>0.995</b> | 64        | 0.995            | 50k         | 200 |
| 10  | 0.0001 | 0.99         | 128       | 0.999            | <b>100k</b> | 500 |
| 11  | 0.001  | 0.95         | 64        | 0.99             | 50k         | 500 |
| 12  | 0.0003 | 0.99         | <b>96</b> | 0.997            | <b>75k</b>  | 500 |

- **Runs 01-04:** LR sensitivity,
- **05-07:**  $\epsilon$  decay ,
- **08-09:**  $\gamma$  effects

# Multi-Intersection Performance Results (Part 1)

| Configuration         | Wait (s) | Queue | Thru. |
|-----------------------|----------|-------|-------|
| <i>Baselines</i>      |          |       |       |
| Fixed-Time            | 660,034  | 589   | 429   |
| Actuated              | 618,820  | 750   | 544   |
| <i>Top RL Configs</i> |          |       |       |
| Run 01                | 1,283    | 83    | 565   |
| Run 05                | 11,069   | 203   | 659   |
| Run 10                | 6,430    | 133   | 682   |
| Run 02                | 18,940   | 254   | 786   |
| Run 04                | 15,221   | 227   | 514   |

## Best: Run 01

- **99.8%** wait reduction
- **85.9%** queue reduction
- LR=0.0005,  $\gamma=0.99$ ,  
 $\epsilon_{decay}=0.995$

## Key Findings

- Slow  $\epsilon$  decay critical
- Moderate LR (0.0001-0.0005)
- 100× better than baseline

# Failed Configurations & Lessons

| Config                     | Wait (s) | Improvement |
|----------------------------|----------|-------------|
| Fixed-Time                 | 660,034  | -           |
| Run 06 ( $\epsilon=0.99$ ) | 443,071  | 32.9%       |
| Run 09 ( $\gamma=0.995$ )  | 357,916  | 45.8%       |
| Run 08 ( $\gamma=0.95$ )   | 161,072  | 75.6%       |
| Run 03 (LR=0.001)          | 62,146   | 90.6%       |
| Run 07 ( $\epsilon=0.98$ ) | 25,707   | 96.1%       |

## Critical Failures

- **Run 06:** Fast  $\epsilon$  (0.99)
  - Premature exploitation
  - Only 33% improvement
- **Run 09:** High  $\gamma$  (0.995)
  - Unstable value estimates
  - Only 46% improvement

## Key Lessons

- Fast  $\epsilon$  decay catastrophic
- Extreme  $\gamma$  hurts
- Proper tuning essential

# Critical: Epsilon Decay Dominates Performance

## Impact

| Decay | Run    | Reduction |
|-------|--------|-----------|
| 0.999 | 05, 10 | 98-99%    |
| 0.995 | 01-04  | 90-99.8%  |
| 0.99  | 06     | 32.9%     |

## Why Critical?

- Delayed feedback through network
- Multi-modal policies
- $8^3 = 512$  joint actions to explore

## Exploration Timeline

Slow (0.999) ✓

Ep 100:  $\epsilon=0.90$ , Ep 500:  $\epsilon=0.61$   
Result: 98-99% improvement

Fast (0.99) ✗

Ep 100:  $\epsilon=0.37$ , Ep 500:  $\epsilon=0.01$   
Result: Only 33% improvement

Use  $\epsilon$  decay  $\geq 0.995$

# Discount Factor: Non-Monotonic Behavior

| $\gamma$    | Horizon     | Result          |
|-------------|-------------|-----------------|
| 0.95        | 100s        | 75.6%           |
| <b>0.99</b> | <b>500s</b> | <b>95-99.8%</b> |
| 0.995       | 1000s       | 45.8%           |

## Too Low (0.95)

- Myopic, reactive only
- Misses future congestion
- 75.6% (mediocre)

## Optimal (0.99)

- Balances near/far future
- Captures traffic waves
- 95-99.8% (excellent)

## Too High (0.995)

- Excessive long-term focus
- Bootstrapping errors accumulate
- Credit assignment fails
- Training unstable
- Only 45.8% (failure!)

Use  $\gamma = 0.99$

# Learning Rate & Batch Size Interactions

| LR            | Run           | Result          |
|---------------|---------------|-----------------|
| 0.0001        | 02, 10        | 97-99%          |
| <b>0.0005</b> | <b>01, 05</b> | <b>98-99.8%</b> |
| 0.001         | 03            | 90%             |
| 0.002         | 04            | 97.7%           |

High LR + Large Batch

Run 04: LR=0.002, Batch=128  
97.7% - *batch stabilizes*

## Insights

- Conservative (0.0001): Stable, slower
- Moderate (0.0005): Best balance
- Aggressive: Needs large batch

High LR + Wrong  $\gamma$

Run 03: LR=0.001, Batch=128  
90.6% - *interactions matter*

## Batch Findings

- 64: Best for moderate LR
- 128: Needed for high LR

**Use LR=0.0005, Batch=64**

# Critical: Early Stopping Essential

## Best vs Final (Run 01)

| Model     | Ep  | Wait (s) |
|-----------|-----|----------|
| Best      | 97  | 1,283    |
| Final     | 500 | 214,763  |
| Degradate |     | × 167!   |

| Run | Best Ep |
|-----|---------|
| 01  | 97      |
| 02  | 186     |
| 05  | 167     |
| 10  | 27      |

## What Failed?

- Training past optimal point
- Overfits to training patterns
- Best at episodes 27-186
- Never at final episode!

## Why Needed?

- Memorizes training seed
- Poor generalization
- Validation plateaus early

# Performance Comparison with Baselines

## Baseline Controllers

### ■ Fixed-Time Control:

- Pre-programmed phase durations
- Through/right: 30s green
- Left turn: 15s green
- Total cycle: 110 seconds

### ■ Actuated Control:

- Vehicle-actuated phases
- Min green: 5s, Max: 60s
- Extension: 3s per detection
- Gap threshold: 3.0s

### ■ Dueling Double DQN (Ours):

- Learned adaptive policy
- Network-wide coordination
- Optimizes cumulative reward

## Multi-Intersection Results

| Method             | Wait (s)     | Throughput    |
|--------------------|--------------|---------------|
| Fixed-Time         | 660,034      | 429           |
| Actuated           | 618,820      | 544           |
| <b>DQN Best</b>    | <b>1,283</b> | <b>565</b>    |
| <b>Improvement</b> | <b>99.8%</b> | <b>+31.7%</b> |

## Single Intersection:

| Method             | Wait (s)      | Queue         |
|--------------------|---------------|---------------|
| RL-DQN             | 488           | 75            |
| Fixed-Time         | 3,523         | 101           |
| Actuated           | 5,656         | 143           |
| <b>Improvement</b> | <b>86-91%</b> | <b>26-48%</b> |

# Key Experimental Findings

## ■ Exploration is Critical:

- Slow  $\epsilon$  decay ( $\geq 0.995$ ): 95-99.8% wait reduction
- Fast decay (0.98-0.99): Only 18-33% reduction
- Extended exploration essential for multi-intersection coordination

## ■ Discount Factor Sensitivity:

- $\gamma = 0.95$  (myopic): 75.6% reduction
- $\gamma = 0.99$  (optimal): 95-99.8% reduction
- $\gamma = 0.995$  (unstable): Only 45.8% reduction

## ■ Early Stopping Matters:

- Best checkpoints: Episodes 27-186
- Final models (Ep 500): Up to 167 $\times$  worse!
- Overfitting to training distribution

## ■ Multi-Intersection Coordination:

- Centralized agent learns implicit coordination
- Best config: 65.5% queue reduction, 53.6% throughput gain
- No explicit communication needed

# Trained Policy in Action

## Policy Demonstration Video

[Click here to watch the video](#)

Trained DQN agent controlling multi-intersection traffic in SUMO

# Summary & Key Findings

- **Problem:** Optimized multi-intersection traffic signal control using Deep RL
- **Approach:**
  - Formulated as MDP with high-dimensional state space (117D)
  - Implemented Dueling Double DQN for stable learning
  - Single agent controls network-wide coordination
- **Results:**
  - **98.3%** reduction in waiting time vs fixed-time (11,069s vs 660,034s)
  - **65.5%** reduction in queue length (203 vs 589 vehicles)
  - **53.6%** improvement in throughput (659 vs 429 vehicles)
  - Stable convergence with proper hyperparameters
- **Key Insights:**
  - Slow  $\epsilon$  decay ( $\geq 0.995$ ) is critical for success
  - Early stopping prevents overfitting ( $167 \times$  degradation observed)
  - Optimal  $\gamma=0.99$  balances myopic vs unstable learning
  - 12-config ablation revealed hyperparameter interactions

# Questions?

---

Thank you for your attention!

Abdul Basit & M. Abdullah  
LUMS